INF8953DE
REINFORCEMENT LEARNING

PRESENTATION

HOW TO COMBINE TREE-SEARCH METHODS IN REINFORCEMENT LEARNING

A REPRODUCIBILITY STUDY

Adam Prévost (1947205)

Alexandre Morinvil (1897222)

Maude Nguyen-The (1843896)

Introduction

- Project track
  - Reproducibility study
  - Yonathan Efroni, Gal Dalal, Bruno Scherrer, and Shie Mannor. How to combine tree-search methods in reinforcement learning, 2019. https://arxiv.org/abs/1809.01843
- Goal of the project
  - Reproduce & validate the paper's result
  - Confirm the mathematical results empirically
  - Compare the results

Presentation plan

1. Introduction
2. **Presentation plan**
3. Paper summary
4. Experiments
5. Conclusion
6. Q & A

Presentation plan

1. Introduction
2. Presentation plan
3. **Paper summary**
4. Experiments
5. Conclusion
6. Q & A

Paper summary

- Problem
  - Lookahead tree search methods don’t always contract
  - Introduces h-greedy consistency to prove contraction
- Proposed algorithms
  - hm-PI
  - hλ-PI (no empirical results)

Presentation plan

1. Introduction
2. Presentation plan
3. Paper summary
4. **Experiments**
5. Conclusion
6. Q & A

Experiments: Environement

- 10 by 10 Gridworld
- One state with reward 1 & the rest with rewards drawn from [-0.1, 0.1]
- Actions : STAY, UP, DOWN, LEFT, RIGHT

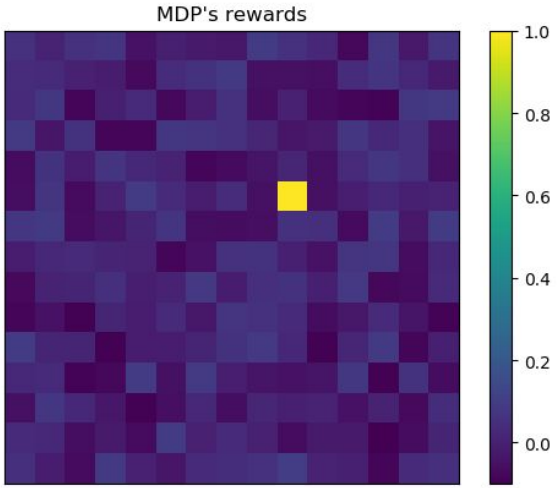

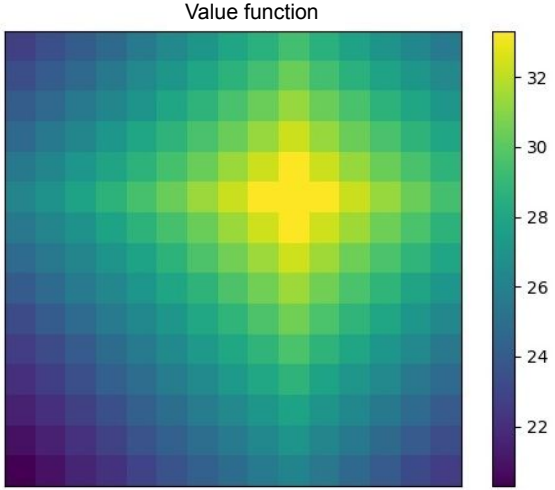

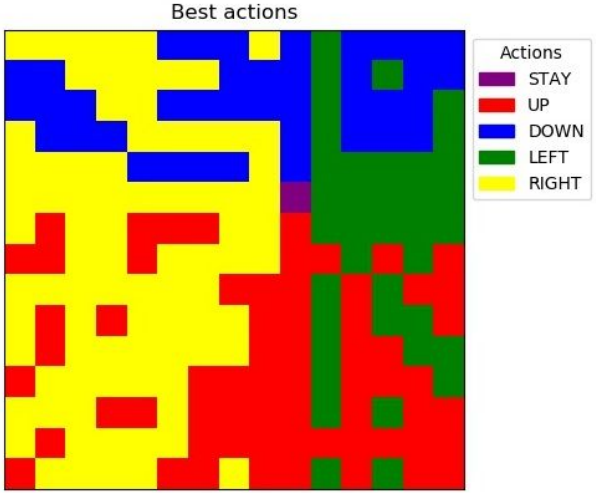

*For example purposes, the figures above are 15 by 15

Experiments: Reproducing hm-PI experiment

- NC-hm-PI & hm-PI using the same *h* and *m* hyper parameters
- Total queries to measure the time performance

Paper's results

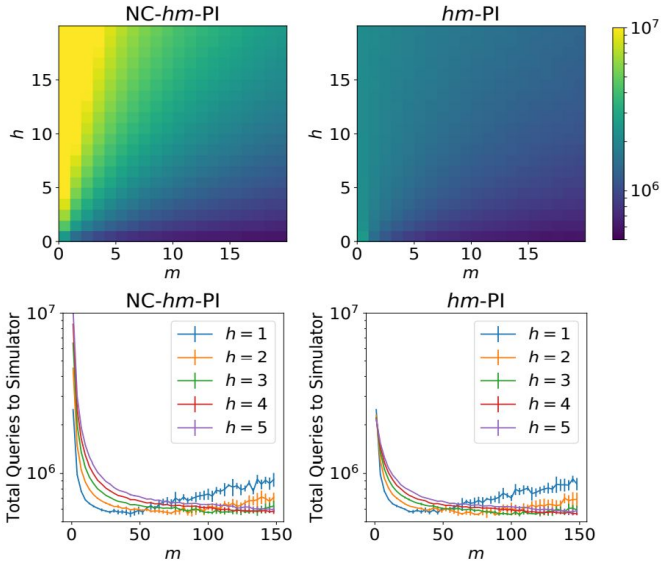

Experimental results

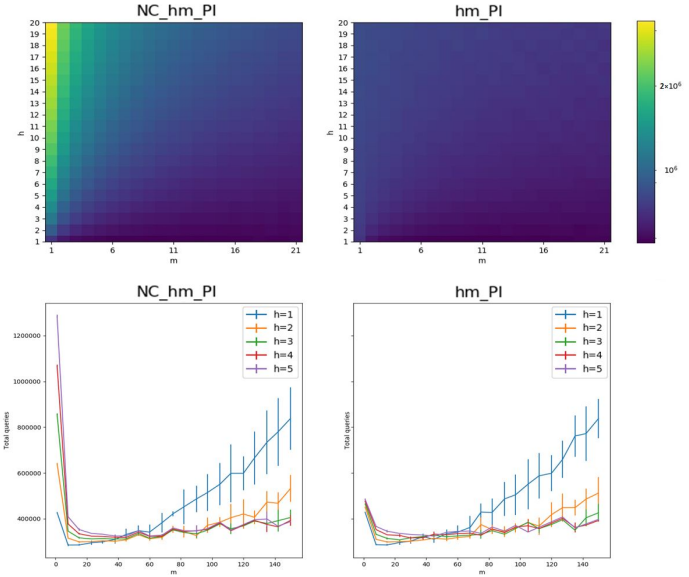

Alexandre Morinvil

Experiments: Reproducing hm-PI experiments

- NC-hm-PI & hm-PI using the same *h* and *m* hyper parameters
- Distance from optimal value function

Paper's results

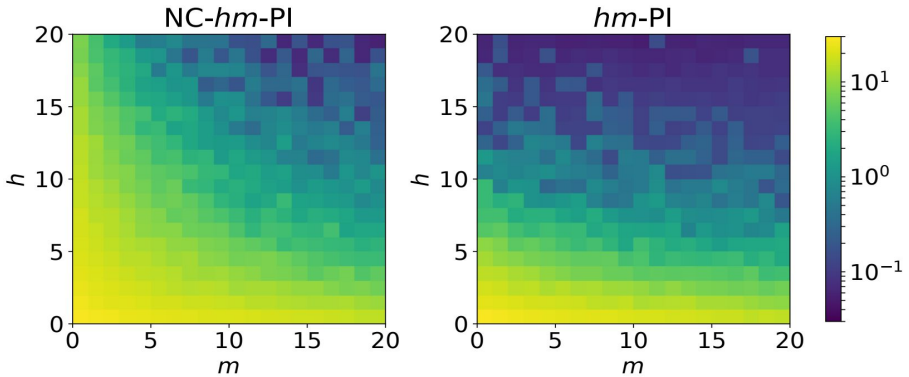

Experimental results

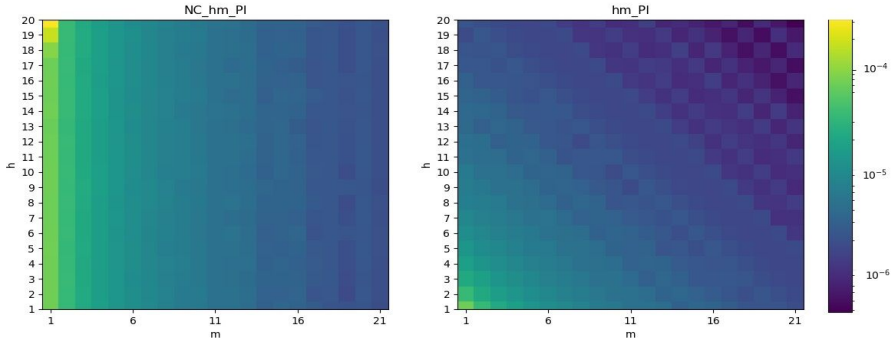

Experiments: Exploring results with hλ-PI

- NC-hλ-PI & hλ-PI using the same *h* and *λ* hyper parameters
- Total queries to measure the time performance

$$T_\lambda^\pi v \stackrel{\text{def}}{=} (1-\lambda)\sum_{j=0}^{\infty}\lambda^j (T^\pi)^{j+1} v$$
$$= v + (I - \gamma\lambda P^\pi)^{-1}(T^\pi v - v).$$

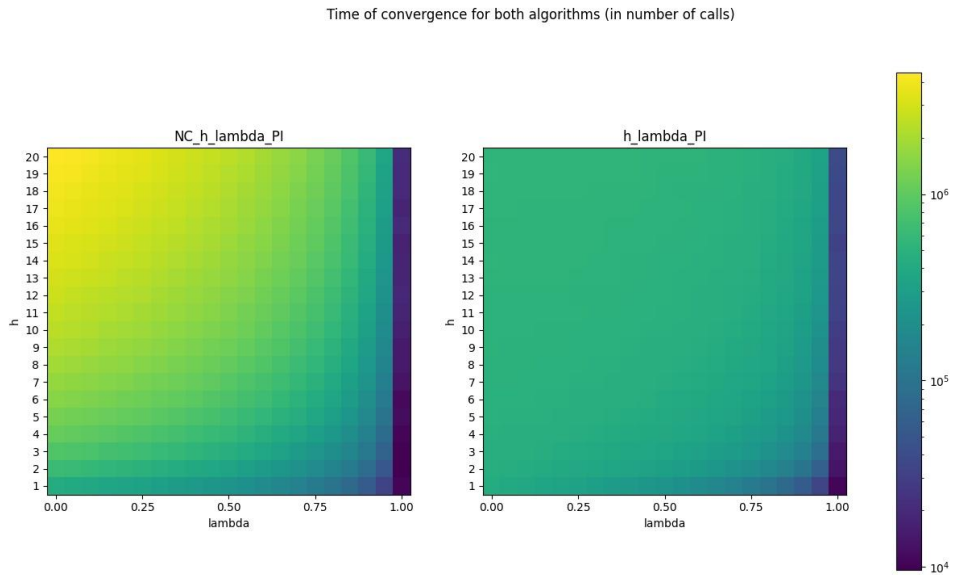

Experiments: Exploring results with hλ-PI

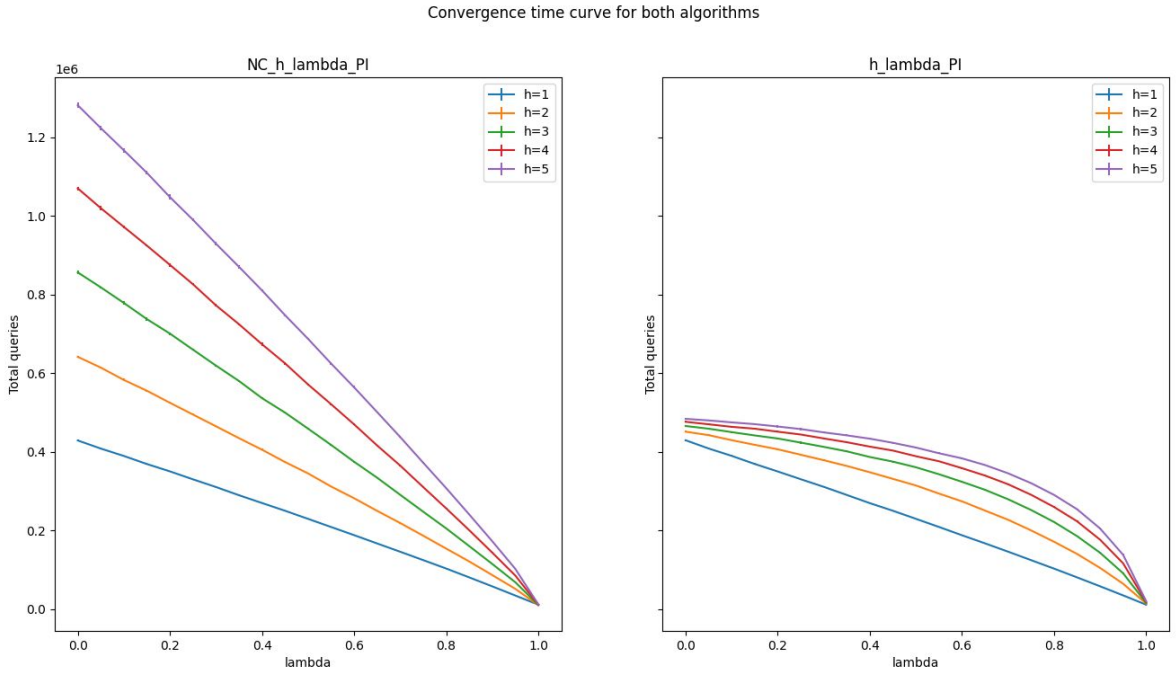

Experiments: Exploring results with hλ-PI

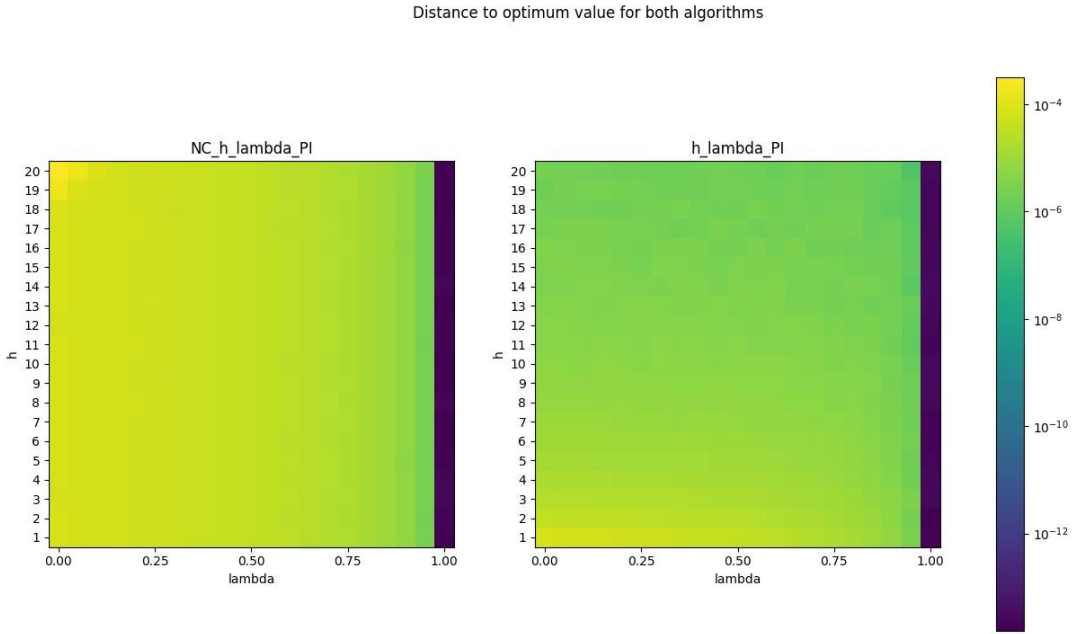

Presentation plan

1. Introduction
2. Presentation plan
3. Paper summary
4. Experiments
5. **Conclusion**
6. Q & A

Conclusion

- Results are similar
- We can affirm the findings reported in the original work
- Reproducibility issues & caveats
  - Missing parameters (ex: the size of the gridworld)
  - Unspecified procedures (ex: how to count queries *in practice*)
  - Unintuitive notation
- Link to our study and code
  https://github.com/AdamPrevost/INF8953.git

Presentation plan

1. Introduction
2. Presentation plan
3. Paper summary
4. Experiments
5. Conclusion
6. **Q & A**

Q & A